

eplg9tdan

March 30, 2023

:#Programming Assignment - 2 : Actor-Critic One-step

```
[1]: '''  
    Installing packages for rendering the game on Colab  
    '''  
  
    !pip install gym pyvirtualdisplay > /dev/null 2>&1  
    !apt-get install -y xvfb python-opengl ffmpeg > /dev/null 2>&1  
    !apt-get update > /dev/null 2>&1  
    !apt-get install cmake > /dev/null 2>&1  
    !pip install --upgrade setuptools 2>&1  
    !pip install ez_setup > /dev/null 2>&1  
    !pip install gym[atari] > /dev/null 2>&1  
    !pip install git+https://github.com/tensorflow/docs > /dev/null 2>&1  
    !pip install gym[classic_control]
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-  
wheels/public/simple/  
Requirement already satisfied: setuptools in /usr/local/lib/python3.9/dist-  
packages (67.6.1)  
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-  
wheels/public/simple/  
Requirement already satisfied: gym[classic_control] in  
/usr/local/lib/python3.9/dist-packages (0.25.2)  
Requirement already satisfied: numpy>=1.18.0 in /usr/local/lib/python3.9/dist-  
packages (from gym[classic_control]) (1.22.4)  
Requirement already satisfied: importlib-metadata>=4.8.0 in  
/usr/local/lib/python3.9/dist-packages (from gym[classic_control]) (6.1.0)  
Requirement already satisfied: cloudpickle>=1.2.0 in  
/usr/local/lib/python3.9/dist-packages (from gym[classic_control]) (2.2.1)  
Requirement already satisfied: gym-notices>=0.0.4 in  
/usr/local/lib/python3.9/dist-packages (from gym[classic_control]) (0.0.8)  
Collecting pygame==2.1.0  
  Downloading  
pygame-2.1.0-cp39-cp39-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (18.3 MB)  
18.3/18.3 MB  
33.6 MB/s eta 0:00:00  
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.9/dist-
```

```

packages (from importlib-metadata>=4.8.0->gym[classic_control]) (3.15.0)
Installing collected packages: pygame
  Attempting uninstall: pygame
    Found existing installation: pygame 2.3.0
    Uninstalling pygame-2.3.0:
      Successfully uninstalled pygame-2.3.0
  Successfully installed pygame-2.1.0

```

```
[ ]: tf.config.list_physical_devices('GPU')
```

```
[ ]: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
```

```

[2]: '''
    A bunch of imports, you don't have to worry about these
    '''

import numpy as np
import random
import torch
import torch.nn as nn
import torch.nn.functional as F
from collections import namedtuple, deque
import torch.optim as optim
import datetime
import gym
from gym.wrappers.record_video import RecordVideo
import glob
import io
import base64
import matplotlib.pyplot as plt
from IPython.display import HTML
from pyvirtualdisplay import Display
import tensorflow as tf
from IPython import display as ipythondisplay
from PIL import Image
import tensorflow_probability as tfp

```

## 0.1 Part 2: One-Step Actor-Critic Algorithm

**Actor-Critic methods** learn both a policy  $\pi(a|s; \theta)$  and a state-value function  $v(s; w)$  simultaneously. The policy is referred to as the actor that suggests actions given a state. The estimated value function is referred to as the critic. It evaluates actions taken by the actor based on the given policy. In this exercise, both functions are approximated by feedforward neural networks.

- The policy network is parametrized by  $\theta$  - it takes a state  $s$  as input and outputs the probabilities  $\pi(a|s; \theta) \forall a$
- The value network is parametrized by  $w$  - it takes a state  $s$  as input and outputs a scalar value associated with the state, i.e.,  $v(s; w)$

- The single step TD error can be defined as follows:

$$\delta_t = R_{t+1} + \gamma v(s_{t+1}; w) - v(s_t; w)$$

- The loss function to be minimized at every step ( $L_{tot}^{(t)}$ ) is a summation of two terms, as follows:

$$L_{tot}^{(t)} = L_{actor}^{(t)} + L_{critic}^{(t)}$$

where,

$$L_{actor}^{(t)} = -\log \pi(a_t | s_t; \theta) \delta_t$$

$$L_{critic}^{(t)} = \delta_t^2$$

- **NOTE: Here, weights of the first two hidden layers are shared by the policy and the value network**
  - First two hidden layer sizes: [1024, 512]
  - Output size of policy network: 2 (Softmax activation)
  - Output size of value network: 1 (Linear activation)

### 0.1.1 Initializing Actor-Critic Network

```
[3]: class ActorCriticModel(tf.keras.Model):
    """
    Defining policy and value networkss
    """
    def __init__(self, action_size, n_hidden1=1024, n_hidden2=512):
        super(ActorCriticModel, self).__init__()

        #Hidden Layer 1
        self.fc1 = tf.keras.layers.Dense(n_hidden1, activation='relu')
        #Hidden Layer 2
        self.fc2 = tf.keras.layers.Dense(n_hidden2, activation='relu')

        #Output Layer for policy
        self.pi_out = tf.keras.layers.Dense(action_size, activation='softmax')
        #Output Layer for state-value
        self.v_out = tf.keras.layers.Dense(1)

    def call(self, state):
        """
        Computes policy distribution and state-value for a given state
        """
        layer1 = self.fc1(state)
        layer2 = self.fc2(layer1)

        pi = self.pi_out(layer2)
        v = self.v_out(layer2)

        return pi, v
```

```
/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283:
DeprecationWarning: `should_run_async` will not call `transform_cell`
automatically in the future. Please pass the result to `transformed_cell`
argument and any exception that happen during the transform in
`preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)
```

### 0.1.2 Agent Class

###Task 2a: Write code to compute  $\delta_t$  inside the Agent.learn() function

```
[4]: class Agent:
    """
    Agent class
    """
    def __init__(self, action_size, lr=0.001, gamma=0.99, seed = 85):
        self.gamma = gamma
        self.ac_model = ActorCriticModel(action_size=action_size)
        self.ac_model.compile(tf.keras.optimizers.Adam(learning_rate=lr))
        np.random.seed(seed)

    def sample_action(self, state):
        """
        Given a state, compute the policy distribution over all actions and
        ↪sample one action
        """
        pi,_ = self.ac_model(state)

        action_probabilities = tfp.distributions.Categorical(probs=pi)
        sample = action_probabilities.sample()

        return int(sample.numpy()[0])

    def actor_loss(self, action, pi, delta):
        """
        Compute Actor Loss
        """
        return -tf.math.log(pi[0,action]) * delta

    def critic_loss(self,delta):
        """
        Critic loss aims to minimize TD error
        """
        return delta**2

    @tf.function
    def learn(self, state, action, reward, next_state, done):
        """
```

*For a given transition (s,a,s',r) update the parameters by computing the gradient of the total loss*

```

"""
with tf.GradientTape(persistent=True) as tape:
    pi, V_s = self.ac_model(state)
    _, V_s_next = self.ac_model(next_state)

    V_s = tf.squeeze(V_s)
    V_s_next = tf.squeeze(V_s_next)

    ##### TO DO: Write the equation for delta (TD error)
    ## Write code below
    delta = reward + self.gamma * V_s_next - V_s ## Complete this
    loss_a = self.actor_loss(action, pi, delta)
    loss_c = self.critic_loss(delta)
    loss_total = loss_a + loss_c

    gradient = tape.gradient(loss_total, self.ac_model.trainable_variables)
    self.ac_model.optimizer.apply_gradients(zip(gradient, self.ac_model.
↪trainable_variables))

```

```
[6]: env = gym.make('Acrobot-v1')
```

```

#Initializing Agent
agent = Agent(lr=1e-4, action_size=env.action_space.n)
#Number of episodes
episodes = 1800
tf.compat.v1.reset_default_graph()

reward_list = []
average_reward_list = []

begin_time = datetime.datetime.now()

steps_his=[]
i=0

for ep in range(1, episodes + 1):
    state = env.reset().reshape(1,-1)
    done = False
    ep_rew = 0
    while not done:
        action = agent.sample_action(state) ##Sample Action
        next_state, reward, done, info = env.step(action) ##Take action
        next_state = next_state.reshape(1,-1)
        ep_rew += reward ##Updating episode reward

```

```

        agent.learn(state, action, reward, next_state, done) ##Update Parameters
        state = next_state ##Updating State
        i+=1
        reward_list.append(ep_rew)
        steps_his.append(i)
        average_reward_list.append(np.mean(reward_list[-10:]))

    if ep % 10 == 0:
        avg_rew = np.mean(reward_list[-10:])
        print('Episode ', ep, 'Reward %f' % ep_rew, 'Average Reward %f' %_
↪avg_rew)

    if ep % 100:
        avg_100 = np.mean(reward_list[-100:])
        if avg_100 > -200:
            print('Stopped at Episode ', ep-100)
            break

time_taken = datetime.datetime.now() - begin_time
print(time_taken)

```

```

/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283:
DeprecationWarning: `should_run_async` will not call `transform_cell`
automatically in the future. Please pass the result to `transformed_cell`
argument and any exception that happen during thetransform in
`preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)
/usr/local/lib/python3.9/dist-packages/gym/core.py:317: DeprecationWarning:
WARN: Initializing wrapper in old step API which returns one bool instead
of two. It is recommended to set `new_step_api=True` to use new step API. This
will be the default behaviour in future.
    deprecation(
/usr/local/lib/python3.9/dist-
packages/gym/wrappers/step_api_compatibility.py:39: DeprecationWarning:
WARN: Initializing environment in old step API which returns one bool
instead of two. It is recommended to set `new_step_api=True` to use new step
API. This will be the default behaviour in future.
    deprecation(

Episode   10 Reward -138.000000 Average Reward -198.700000
Stopped at Episode   -90
0:00:25.487957

```

```


[7]: plt.style.use('classic')
     plt.figure(figsize=(14,9))

```

```

plt.plot(np.arange(len(reward_list)),reward_list,label='Point episode_
↪scores',color='yellow')
plt.plot(np.arange(len(average_reward_list)),average_reward_list,label='Running_
↪average over 10 episodes',color='red')
plt.xlabel('Episodes',fontsize=20)
plt.ylabel('Rewards',fontsize=20)
plt.title('Reward curve for Acrobot one step return configuration',fontsize=20)
plt.legend(loc='lower right')
plt.figure(figsize=(14,9))
plt.plot(np.arange(len(steps_his)),steps_his,label='Steps in each episode')
plt.xlabel('Episodes',fontsize=20)
plt.ylabel('Steps',fontsize=20)
plt.title('Steps in every episode for Acrobot one step return_
↪configuration',fontsize=20)
plt.legend(loc='lower right')
plt.show()

```





```
[8]: env = gym.make('CartPole-v1')

#Initializing Agent
agent = Agent(lr=1e-4, action_size=env.action_space.n)
#Number of episodes
episodes = 1800
tf.compat.v1.reset_default_graph()

reward_list = []
average_reward_list = []
begin_time = datetime.datetime.now()

steps_his=[]
i=0

for ep in range(1, episodes + 1):
    state = env.reset().reshape(1,-1)
    done = False
    ep_rew = 0
    i=0
    while not done:
        action = agent.sample_action(state) ##Sample Action
```



```

        next_state, reward, done, info = env.step(action) ##Take action
        next_state = next_state.reshape(1,-1)
        ep_rew += reward ##Updating episode reward
        agent.learn(state, action, reward, next_state, done) ##Update Parameters
        state = next_state ##Updating State
        i+=1
    reward_list.append(ep_rew)
    steps_his.append(i)
    average_reward_list.append(np.mean(reward_list[-10:]))

    if ep % 10 == 0:
        avg_rew = np.mean(reward_list[-10:])
        print('Episode ', ep, 'Reward %f' % ep_rew, 'Average Reward %f' %
↪avg_rew)

    if ep % 100:
        avg_100 = np.mean(reward_list[-100:])
        if avg_100 > 195.0:
            print('Stopped at Episode ', ep-100)
            break

time_taken = datetime.datetime.now() - begin_time
print(time_taken)

```

```

/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283:
DeprecationWarning: `should_run_async` will not call `transform_cell`
automatically in the future. Please pass the result to `transformed_cell`
argument and any exception that happen during thetransform in
`preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)
/usr/local/lib/python3.9/dist-packages/gym/core.py:317: DeprecationWarning:
WARN: Initializing wrapper in old step API which returns one bool instead
of two. It is recommended to set `new_step_api=True` to use new step API. This
will be the default behaviour in future.
    deprecation(
/usr/local/lib/python3.9/dist-
packages/gym/wrappers/step_api_compatibility.py:39: DeprecationWarning:
WARN: Initializing environment in old step API which returns one bool
instead of two. It is recommended to set `new_step_api=True` to use new step
API. This will be the default behaviour in future.
    deprecation(
Episode 10 Reward 34.000000 Average Reward 34.900000
Episode 20 Reward 80.000000 Average Reward 61.200000
Episode 30 Reward 56.000000 Average Reward 51.800000
Episode 40 Reward 47.000000 Average Reward 72.100000

```

```

Episode 50 Reward 146.000000 Average Reward 85.500000
Episode 60 Reward 73.000000 Average Reward 93.800000
Episode 70 Reward 120.000000 Average Reward 85.800000
Episode 80 Reward 72.000000 Average Reward 80.100000
Episode 90 Reward 68.000000 Average Reward 77.600000
Episode 100 Reward 126.000000 Average Reward 101.600000
Episode 110 Reward 66.000000 Average Reward 115.400000
Episode 120 Reward 103.000000 Average Reward 81.400000
Episode 130 Reward 44.000000 Average Reward 57.700000
Episode 140 Reward 59.000000 Average Reward 47.300000
Episode 150 Reward 104.000000 Average Reward 69.700000
Episode 160 Reward 92.000000 Average Reward 59.800000
Episode 170 Reward 89.000000 Average Reward 117.400000
Episode 180 Reward 90.000000 Average Reward 105.500000
Episode 190 Reward 118.000000 Average Reward 114.200000
Episode 200 Reward 126.000000 Average Reward 119.600000
Episode 210 Reward 63.000000 Average Reward 115.500000
Episode 220 Reward 89.000000 Average Reward 95.000000
Episode 230 Reward 177.000000 Average Reward 93.100000
Episode 240 Reward 68.000000 Average Reward 99.700000
Episode 250 Reward 61.000000 Average Reward 95.800000
Episode 260 Reward 231.000000 Average Reward 145.100000
Episode 270 Reward 152.000000 Average Reward 172.400000
Episode 280 Reward 225.000000 Average Reward 226.100000
Episode 290 Reward 209.000000 Average Reward 387.600000
Episode 300 Reward 113.000000 Average Reward 274.900000
Episode 310 Reward 125.000000 Average Reward 113.600000
Episode 320 Reward 191.000000 Average Reward 162.700000
Episode 330 Reward 218.000000 Average Reward 191.900000
Stopped at Episode 233
0:05:53.312282


```

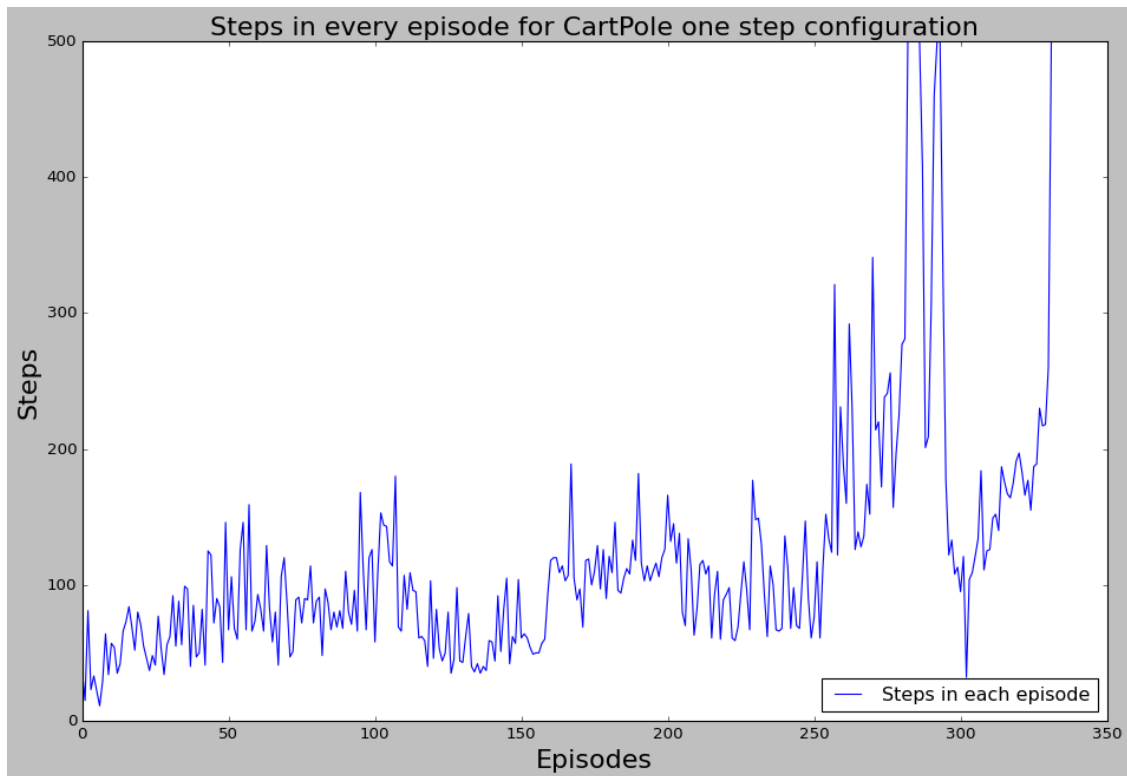
```

[9]: plt.style.use('classic')
plt.figure(figsize=(14,9))
plt.plot(np.arange(len(reward_list)),reward_list,label='Point episode_
↪scores',color='yellow')
plt.plot(np.arange(len(average_reward_list)),average_reward_list,label='Running_
↪average over 10 episodes',color='red')
plt.xlabel('Episodes',fontsize=20)
plt.ylabel('Rewards',fontsize=20)
plt.title('Reward curve for CartPole one step configuration',fontsize=20)
plt.legend(loc='lower right')
plt.figure(figsize=(14,9))
plt.plot(np.arange(len(steps_his)),steps_his,label='Steps in each episode')
plt.xlabel('Episodes',fontsize=20)
plt.ylabel('Steps',fontsize=20)

```

```
plt.title('Steps in every episode for CartPole one step_↵  
↵configuration',fontsize=20)  
plt.legend(loc='lower right')  
plt.show()
```





```
[10]: env = gym.make('MountainCar-v0')

#Initializing Agent
agent = Agent(lr=1e-4, action_size=env.action_space.n)
#Number of episodes
episodes = 1800
tf.compat.v1.reset_default_graph()

reward_list = []
average_reward_list = []
begin_time = datetime.datetime.now()

steps_his=[]
i=0

for ep in range(1, episodes + 1):
    state = env.reset().reshape(1,-1)
    done = False
    ep_rew = 0
    i=0
    while not done:
        action = agent.sample_action(state) ##Sample Action
```

```

        next_state, reward, done, info = env.step(action) ##Take action
        next_state = next_state.reshape(1,-1)
        ep_rew += reward ##Updating episode reward
        agent.learn(state, action, reward, next_state, done) ##Update Parameters
        state = next_state ##Updating State
        i+=1
    reward_list.append(ep_rew)
    steps_his.append(i)
    average_reward_list.append(np.mean(reward_list[-10:]))

    if ep % 10 == 0:
        avg_rew = np.mean(reward_list[-10:])
        print('Episode ', ep, 'Reward %f' % ep_rew, 'Average Reward %f' %
↪avg_rew)

    if ep % 100:
        avg_100 = np.mean(reward_list[-100:])
        if avg_100 > -150:
            print('Stopped at Episode ', ep-100)
            break

time_taken = datetime.datetime.now() - begin_time
print(time_taken)

```

```

/usr/local/lib/python3.9/dist-packages/ipykernel/ipkernel.py:283:
DeprecationWarning: `should_run_async` will not call `transform_cell`
automatically in the future. Please pass the result to `transformed_cell`
argument and any exception that happen during thetransform in
`preprocessing_exc_tuple` in IPython 7.17 and above.
    and should_run_async(code)
/usr/local/lib/python3.9/dist-packages/gym/core.py:317: DeprecationWarning:
WARN: Initializing wrapper in old step API which returns one bool instead
of two. It is recommended to set `new_step_api=True` to use new step API. This
will be the default behaviour in future.
    deprecation(
/usr/local/lib/python3.9/dist-
packages/gym/wrappers/step_api_compatibility.py:39: DeprecationWarning:
WARN: Initializing environment in old step API which returns one bool
instead of two. It is recommended to set `new_step_api=True` to use new step
API. This will be the default behaviour in future.
    deprecation(
Episode 10 Reward -200.000000 Average Reward -200.000000
Episode 20 Reward -200.000000 Average Reward -200.000000
Episode 30 Reward -200.000000 Average Reward -200.000000
Episode 40 Reward -200.000000 Average Reward -200.000000

```

[illegible]

[illegible]

[illegible]



```

Episode 1490 Reward -200.000000 Average Reward -200.000000
Episode 1500 Reward -200.000000 Average Reward -200.000000
Episode 1510 Reward -200.000000 Average Reward -200.000000
Episode 1520 Reward -200.000000 Average Reward -200.000000
Episode 1530 Reward -200.000000 Average Reward -200.000000
Episode 1540 Reward -200.000000 Average Reward -200.000000
Episode 1550 Reward -200.000000 Average Reward -200.000000
Episode 1560 Reward -200.000000 Average Reward -200.000000
Episode 1570 Reward -200.000000 Average Reward -200.000000
Episode 1580 Reward -200.000000 Average Reward -200.000000
Episode 1590 Reward -200.000000 Average Reward -200.000000
Episode 1600 Reward -200.000000 Average Reward -200.000000
Episode 1610 Reward -200.000000 Average Reward -200.000000
Episode 1620 Reward -200.000000 Average Reward -200.000000
Episode 1630 Reward -200.000000 Average Reward -200.000000
Episode 1640 Reward -200.000000 Average Reward -200.000000
Episode 1650 Reward -200.000000 Average Reward -200.000000
Episode 1660 Reward -200.000000 Average Reward -200.000000
Episode 1670 Reward -200.000000 Average Reward -200.000000
Episode 1680 Reward -200.000000 Average Reward -200.000000
Episode 1690 Reward -200.000000 Average Reward -200.000000
Episode 1700 Reward -200.000000 Average Reward -200.000000
Episode 1710 Reward -200.000000 Average Reward -200.000000
Episode 1720 Reward -200.000000 Average Reward -200.000000
Episode 1730 Reward -200.000000 Average Reward -200.000000
Episode 1740 Reward -200.000000 Average Reward -200.000000
Episode 1750 Reward -200.000000 Average Reward -200.000000
Episode 1760 Reward -200.000000 Average Reward -200.000000
Episode 1770 Reward -200.000000 Average Reward -200.000000
Episode 1780 Reward -200.000000 Average Reward -200.000000
Episode 1790 Reward -200.000000 Average Reward -200.000000
Episode 1800 Reward -200.000000 Average Reward -200.000000
0:54:06.593545


```


```

[11]: plt.style.use('classic')
plt.figure(figsize=(14,9))
plt.plot(np.arange(len(reward_list)),reward_list,label='Point episode_
↪scores',color='yellow')
plt.plot(np.arange(len(average_reward_list)),average_reward_list,label='Running_
↪average over 10 episodes',color='red')
plt.xlabel('Episodes',fontsize=20)
plt.ylabel('Rewards',fontsize=20)
plt.title('Reward curve for MountainCar one step return_
↪configuration',fontsize=20)
plt.legend(loc='lower right')
plt.figure(figsize=(14,9))
plt.plot(np.arange(len(steps_his)),steps_his,label='Steps in each episode')

```

```
plt.xlabel('Episodes',fontsize=20)
plt.ylabel('Steps',fontsize=20)
plt.title('Steps in every episode for MountainCar one step return_
↳configuration',fontsize=20)
plt.legend(loc='lower right')
plt.show()
```





### 0.1.3 Code for rendering ([source](#))

```
[ ]: # Render an episode and save as a GIF file

display = Display(visible=0, size=(400, 300))
display.start()

def render_episode(env: gym.Env, model: tf.keras.Model, max_steps: int):
    screen = env.render(mode='rgb_array')
    im = Image.fromarray(screen)

    images = [im]

    state = tf.constant(env.reset(), dtype=tf.float32)
    for i in range(1, max_steps + 1):
        state = tf.expand_dims(state, 0)
        action_probs, _ = model(state)
        action = np.argmax(np.squeeze(action_probs))
        state, _, done, _ = env.step(action)
        state = tf.constant(state, dtype=tf.float32)
```

```

# Render screen every 10 steps
if i % 10 == 0:
    screen = env.render(mode='rgb_array')
    images.append(Image.fromarray(screen))

if done:
    break

return images

# Save GIF image
images = render_episode(env, agent.ac_model, 200)
image_file = 'cartpole-v1.gif'
# loop=0: loop forever, duration=1: play each frame for 1ms
images[0].save(
    image_file, save_all=True, append_images=images[1:], loop=0, duration=1)

```

/usr/local/lib/python3.8/dist-packages/gym/core.py:43: DeprecationWarning:  
 WARN: The argument mode in render method is deprecated; use render\_mode  
 during environment initialization instead.

See here for more information: <https://www.gymnasium.ml/content/api/deprecation>

```

[ ]: import tensorflow_docs.vis.embed as embed
     embed.embed_file(image_file)

```

```

[ ]: <IPython.core.display.HTML object>

```